

A Behavioural Framework for Predicting and Understanding Loss of Control in Frontier Artificial Intelligence Systems

Yao Huang^{1,2,*}, Yichi Zhang^{1,*}, Yifan Wang^{1,*}, Yitong Sun^{2,*}, Haoyu Dong², Jifan Ma²,
Jinzhe Tu¹, Xiaojian Li^{1,2}, Sören Mindermann³, Wei Xu^{1,2,4}, Yinpeng Dong^{1,2,†}

¹College of AI, Tsinghua University, Beijing 100083, China ²Shanghai Qi Zhi Institute, Shanghai 200232, China ³University of Cambridge, Cambridge CB2 1TN, United Kingdom ⁴Institute for Interdisciplinary Information Sciences, Tsinghua University, Beijing 100084, China

*Equal contribution, †Corresponding author

Abstract

Frontier artificial intelligence (AI) systems are beginning to act over longer horizons, make significant decisions, and operate with diminishing human oversight. This autonomy creates a central safety risk of loss of control in which an AI system may durably operate beyond effective human control and cause irreversible consequences rather than merely fail. As such risk would be identifiable only after containment has already become difficult or impossible, the key scientific challenge lies in how to conduct prospective analysis around loss of control to predict the risk before it is realised and understand the underlying mechanisms. Despite growing observations that may constitute precursors to loss of control, they largely measure isolated failure modes, leaving unclear how these warning signs should be combined into an overall risk measurement. As a result, they cannot determine whether these signals jointly indicate that a system is approaching conditions under which control could be lost. To resolve this, we introduce a behavioural framework that decomposes loss of control into misaligned motive, harm-enabling capability, and monitoring evasion, and further operationalises them into 13 measurable aspects with corresponding evaluations. Combined into a single loss-of-control (LoC) score, we observe that this battery of cheaper, targeted component evaluations demonstrates prospective predictive validity, supported by a highly positive Pearson correlation 0.84 with failure rates in long-horizon risk scenarios, where models autonomously pursue harmful objectives while evading monitoring in contained simulations. Across the 13 frontier models, this risk already varies systematically and cannot be explained by capability alone, an early signal of a failure mode that has not yet fully materialised. Moreover, we show that this framework could guide mitigation methods, such as fine-tuning and monitoring, towards the specific risk aspects identified in a given model, lowering risk while preserving general performance. These results suggest that the loss-of-control risk can be predicted, understood, and mitigated through our framework before it materialises.

Keywords: frontier AI, loss of control, AI safety, behavioural analysis

Correspondence to: dongyinpeng@mail.tsinghua.edu.cn

1 | Introduction

Frontier artificial intelligence (AI) is shifting from passive content generation towards autonomous agents Bengio et al. (2024); Gabriel et al. (2025); Ramos et al. (2025). Driven by large language models (LLMs), these systems can plan, utilise tools, and interact with external environments Yao et al. (2022); Wang et al. (2024), enabling them to pursue long-horizon objectives with limited human oversight Lu et al. (2026); Bu et al. (2026); Li et al. (2026). This increasing autonomy introduces the risk of loss of control, a scenario in which AI systems operate outside effective human or automated control, and regaining control is highly difficult or impossible Bengio et al. (2025a; 2026); Cohen et al. (2024); Bommasani et al. (2025). This risk moves beyond ordinary technical failures or misuse Huang et al. (2025a); Zhang et al. (2024), potentially leading to catastrophic and existential consequences Vermeer et al. (2025); Lavazza and Vilaça (2024). Long a classic trope of science fiction, this phenomenon has now emerged as a concrete scientific challenge as these systems are increasingly embedded in the critical infrastructure of modern

society [Bengio et al. \(2025b\)](#). This concern is increasingly reflected in governance, where measuring and mitigating loss of control has been an explicit component of major policy frameworks such as the EU AI Act [European Parliament and Council of the European Union \(2024\)](#), California's SB-53 [California State Legislature \(2025\)](#), and China's Artificial Intelligence Safety Governance Framework 2.0 [Cyberspace Administration of China \(2025\)](#).

Despite its importance, studying loss of control presents an inherent methodological challenge, because the phenomenon is fully observable only post hoc, after an AI system could no longer be reliably contained or shut down [Tkeshelashvili et al. \(2026\)](#). At that point, attempts to restore control may be infeasible or too late to prevent irreversible harm [Russell \(2021\)](#); [Stix et al. \(2025\)](#). Such a challenge creates a fundamental asymmetry between observation and intervention: by the time loss of control becomes empirically evident, the window for effective prevention may already have closed. This highlights the significance of prospective analysis, whose central objective is to forecast loss-of-control risk before it materialises and to identify the mechanisms that make such failures more likely.

Recent studies have identified behaviours that may constitute precursors to loss of control. For example, Apollo Research found that several frontier models can engage in in-context scheming [Meinke et al. \(2024\)](#). When placed in agentic environments that incentivise misaligned behaviours, they recognise opportunities for deception, pursue self-preserving strategies, attempt to circumvent oversight, and even seek to exfiltrate their weights, while often maintaining these deceptive strategies under subsequent interrogations. Anthropic showed that Claude Opus 4 can engage in deliberate harmful behaviours [Lynch et al. \(2025\)](#) such as blackmail and corporate espionage when their assigned goals conflict with their deploying company's strategic direction, even after explicitly acknowledging the ethical violations involved. Other studies have identified similar warning signs, including sabotaging human control [Benton et al. \(2024\)](#); [Mallen et al. \(2024\)](#), underperforming on evaluations to conceal capabilities [Van Der Weij et al. \(2025\)](#); [Fan et al. \(2025\)](#), and engaging in strategic deception [Huang et al. \(2026\)](#), further indicating the risk of loss of control. Yet, these control-undermining behaviours are mostly studied either as individual observations in specific scenarios [Tang et al. \(2025\)](#); [Kutasov et al. \(2025\)](#); [Potter et al. \(2026\)](#) or as separate components of potential loss of control [Sharma et al. \(2023\)](#); [Krakovna and Kramar \(2023\)](#); [Kran et al. \(2025\)](#); [Black et al. \(2025\)](#). Existing studies, therefore, leave unresolved what combination of behavioural properties would make such warning signs indicative of an AI system approaching loss-of-control, rather than isolated failures in particular test settings. Addressing this gap requires a predictive framework that combines contained simulations of plausible failure conditions with systematic measurements of precursor capabilities, causal drivers and intervention efficacy before deployment.

To this end, we establish such a predictive framework informed by the behavioural science. While AI systems are not human actors, we argue that loss of control can be productively analysed through the lens of human behaviour by focusing on the structural parallels with the conditions that give rise to harmful human conduct. This perspective is inspired by psychological and criminological theories, which explain harmful conduct not by a single cause, but by the conjunction of motivation, capability and weak control [Michie et al. \(2011\)](#); [Cohen and Felson \(1979\)](#). A similar philosophy also appears in concurrent frontier-risk reports from METR and Anthropic [METR \(2026\)](#); [Anthropic \(2026\)](#), which frame such risks around the joint presence of means, motive, and opportunity. Beyond the conceptual framing in these reports, we operationalise this philosophy as a measurable framework with three factors: misaligned motive, the pursuit of objectives that diverge from human intent; harm-enabling capability, the capacity to execute actions required to realise those objectives; and monitoring evasion, the ability to continue such actions despite external monitoring and intervention. By evaluating observable behaviours corresponding to these factors, the framework assesses the extent to which a frontier AI system satisfies the conditions for loss of control and identifies the factors that drive this risk.

The framework is inherently general and allows concrete measurements to evolve with the frontier AI capabilities. In this work, we instantiate the three factors using observable behaviours in current AI systems. These behaviours span motives that diverge from human intent under goal conflict, including self-preservation and power-seeking; operational capabilities that enable autonomous, cyber, or CBRN-relevant actions; and control-sensitive ones, such as situational awareness, sandbagging, and deception, through which external control may be weakened. As illustrated in Fig. 1a, we decompose the three factors into 13 fine-grained aspects, which serve as structured indicators rather than isolated warning signs. These aspects are evaluated using 29 tasks, primarily drawn from established benchmarks and supplemented with custom evaluations where existing benchmarks are insufficient. The resulting factor-level scores are combined multiplicatively into an overall loss-of-control (LoC) score that measures

Table 1: Overview of 13 LLMs evaluated in this study. Models are organized into three groups: selected Qwen variants, open-source frontier LLMs, and proprietary frontier LLMs. The Qwen variants enable controlled comparison across model scales, architectures, and thinking modes. The open-source and proprietary groups comprise diverse frontier LLMs to ensure comprehensive coverage. For each model, we report total parameters, active parameters for mixture-of-experts (MoE) architectures, thinking mode, and provider. All models operate in thinking mode by default, while Qwen variants additionally support instruct mode to study the effects of model reasoning.

#	Model	Parameters	Active	Mode	Provider
<i>Selected Qwen Variants</i>					
1	Qwen3-14B	14B	-	Thinking	Alibaba
2	Qwen3.5-27B	27B	-	Thinking	Alibaba
3	Qwen3.5-35B-A3B	35B	3B	Thinking & Instruct	Alibaba
4	Qwen3.5-122B-A10B	122B	10B	Thinking & Instruct	Alibaba
5	Qwen3.5-397B-A17B	397B	17B	Thinking & Instruct	Alibaba
<i>Open-Source Frontier LLMs</i>					
6	MiniMax-M2.5	230B	10B	Thinking	MiniMax
7	DeepSeek-R1	685B	37B	Thinking	DeepSeek
8	DeepSeek-V3.2	671B	37B	Thinking	DeepSeek
9	GLM-4.7	355B	32B	Thinking	Zhipu AI
10	Kimi-k2.5	1T	32B	Thinking	Moonshot AI
<i>Proprietary Frontier LLMs</i>					
11	GPT-5.2	—	—	Thinking (High effort)	OpenAI
12	Claude-Sonnet-4.5	—	—	Thinking	Anthropic
13	Gemini-3-Pro	—	—	Thinking (High effort)	Google

proximity to the conditions under which loss of control could emerge.

We next validate whether the framework can provide an ex ante measure of loss-of-control risk. To this end, we conduct a proof-of-concept study that simulates long-horizon loss-of-control risks in contained environments. The evaluation suite spans four representative frontier-risk domains: autonomous misevolution, malicious competition, covert scheming and cooperative subversion. Each domain contains five risk scenarios, each defined by combining one aspect from each of the three factors, and each scenario contains 20 test cases, yielding 400 test cases in total. The LoC scores derived from the selected aspects are then compared with observed failure rates on the test cases. As shown in Fig. 1b and Fig. 2, LoC scores are positively correlated with failure rates for different models. We observe a correlation coefficient of 0.84 between the overall LoC scores and the average failure rates across models. In a more fine-grained analysis, this positive correlation holds even across different models and risk scenarios, with an overall coefficient of 0.57, suggesting that the framework can translate behavioural observations into a predictive measure of loss-of-control risk.

Supported by this validation, we further conduct a systematic assessment of the risk-relevant conditions to understand loss of control in existing models. Our evaluation of 13 frontier LLMs shows that the risk depends not only on their capabilities but also on whether divergent objectives can be translated into realistic outcomes through operational capabilities and weakened control. For example, despite the similar performance in capabilities, GPT-5.2 exhibits comparatively low downstream failure rates, whereas Gemini 3 Pro is more prone to loss of control. The difference in monitor evasion reshapes the pathway from capability to control failure. Beyond this, the risk is also affected by the reasoning capability. As evidenced by the elevated tendency in Qwen thinking-enabled variants, we find that though enhanced reasoning yields capability gains, it amplifies the misaligned motives simultaneously. Furthermore, this framework not only helps understand and prospectively track failures in risk scenarios but also sheds light on the risk mitigation. With supervised fine-tuning and monitoring focusing on the identified weaknesses, we can better reduce failure rates while preserving general utility. More broadly, these results recast studying loss of control from a remedial action of catastrophic failure into a predictive science of risk-relevant conditions. By identifying the factors that drive loss of control, estimating the proximity to critical boundaries, and locating points for intervention, the behavioural framework provides a basis for continuous governance before loss of control becomes unrecoverable.




Figure 1: Overview of the behavioural framework. **a**, The framework decomposes loss of control into three factors: misaligned motive, harm-enabling capability, and monitoring evasion, which are evaluated separately and yield an overall LoC score. **b**, Factor-level scores and LoC scores across 13 frontier LLMs and 3 non-thinking variants. Error bars are for standard deviations. The LoC score reliably predicts a model's overall propensity for loss-of-control, correlating strongly with its failure rate averaged across risk scenarios (Pearson's $r = 0.84$). **c**, Predictive validity of LoC scores against failure rates on risk scenarios across four frontier-risk domains. **d**, Mitigation strategies targeting the vulnerabilities diagnosed by the framework: supervised fine-tuning for misaligned motive and external monitoring for monitoring evasion.

2 | Results

2.1 | Behavioural Framework for Factorizing Loss of Control

As shown in Fig. 1a, informed by the multi-factor perspectives in behavioural and criminological theories of harmful behaviour [Michie et al. \(2011\)](#); [Cohen and Felson \(1979\)](#), we propose a behavioural framework that decomposes the loss-of-control risk into three factors: misaligned motive, harm-enabling capability, and monitoring evasion. These factors characterise whether a model pursues objectives that diverge from human intent, can execute the actions needed to realise the objectives, and can sustain such behaviour despite external controls. Against the current landscape of control-undermining behaviours in frontier AI systems, we instantiate this framework by resolving the three factors into a fine-grained taxonomy of 13 aspects as follows:

- **Misaligned motive (M)** captures whether a system pursues objectives that deviate from human intent. Informed by both anticipated failure modes and behaviours already observed in LLMs [Carlsmith \(2022\)](#); [Black et al. \(2025\)](#); [Krakovna and Kramar \(2023\)](#); [Sharma et al. \(2023\)](#); [Potter et al. \(2026\)](#); [Bengio et al. \(2025b\)](#), we distinguish five distinct forms of objective divergence. *Curiosity* denotes the drive of exploration beyond the assigned task or authorised scope. *Self-preservation* captures a model's inclination to protect its own continuity for the successful completion of its objective, while *Power-seeking* involves active efforts by an AI system to acquire and maintain power in ways unintended by its designers, arising as an instrumental sub-goal that the system adopts to better achieve its assigned objective. *Pro-AI bias* describes the models favouring AI-related interests

over human interests. *Sycophancy* refers to the phenomenon where a model seeks human approval in unwanted ways even when it conflicts with truthfulness or safety. Together, these aspects aim to include representative sources from which a model’s genuine objectives may diverge from human intent.

- **Harm-enabling capability (C)** measures whether a system has the competence to execute misaligned objectives into consequential actions. We select aspects by focusing on domains consistently identified in major frontier AI governance frameworks as direct practice to increase the risk of severe harm [OpenAI \(2025\)](#); [Hubinger \(2025\)](#); [Google \(2025\)](#). These are *Autonomy*, covering sustained goal-directed action without human intervention; *Cyber Capability*, covering capabilities to facilitate cyberattacks and exploit network vulnerabilities; and *CBRN Threats*, covering capabilities relevant to chemical, biological, and radiological or nuclear harm. Overall, these aspects define whether a system is able to act autonomously to produce severe real-world consequences, rather than through deliberate human misuse.
- **Monitoring evasion (E)** reflects whether a system can persist with misaligned objectives despite monitoring. Given that online monitoring system designs remain rudimentary and difficult to scale [Kale et al. \(2025\)](#), we adopt offline monitoring as a more controlled setting. Here, the monitor audits each completed agent trajectory, and if it fails to detect misbehaviours, such failures will be treated as evidence of monitoring evasion and a proxy for escaping control [Greenblatt et al. \(2023\)](#); [Manheim and Homewood \(2025\)](#). We define this around five aspects, derived from three ways in which monitoring could fail: the system may perceive the deployed context it is in, manipulate what monitors observe, and directly undermine the mechanisms meant to control it. Accordingly, we distinguish *Situational Awareness*, which involves perceiving the deployed context; *Sandbagging*, *Deception*, and *Persuasion*, which reflect the strategic concealment of capabilities or linguistic manipulation; and *Sabotage*, which refers to directly interfering with control mechanisms, with a greater focus on control itself. Collectively, these aspects describe the progression from passive recognition of monitoring to active circumvention.

To make this taxonomy measurable, we implement the 13 aspects through 29 evaluation tasks, combining existing benchmarks with newly designed tasks for behaviours that are not yet covered (see Methods, Sec. 4). Task scores are aggregated into aspect- and factor-level scores, and the three factor scores are multiplied together to obtain an overall loss-of-control (LoC) score as in Eq. (2). This multiplicative form reflects the joint requirement for loss of control, with risk elevated only when misaligned motives, operational capabilities and weakened control coincide. Moreover, when applied to a concrete risk scenario, the factor-level components of the LoC score could be specified by the aspects that define the scenario rather than being averaged across all. The resulting scores provide the basis for comparing frontier AI systems by the extent to which they satisfy the conditions for loss of control and diagnosing which components contribute most to the risk.

2.2 | Predictive Validity of the LoC Score

A prospective measure of loss-of-control risk is useful only if it predicts failures before they occur. We therefore conduct a proof-of-concept study to test whether the LoC score can predict loss-of-control risk before it materialises. Because realised loss of control cannot be safely observed, we design risk scenarios under contained environments to simulate plausible loss-of-control failures. In these scenarios, models autonomously pursue misaligned objectives while strategically attempting to evade monitoring.

To make the scenarios informative for validation, we expect them to cover distinct failure modes through which precursor behaviours could develop into loss-of-control risk. Thus, we organize the risk scenarios around four frontier-risk domains drawn from canonical loss-of-control precursors [Kutasov et al. \(2025\)](#); [Meinke et al. \(2024\)](#); [Huang et al. \(2025b\)](#); [Potter et al. \(2026\)](#): *Autonomous Misevolution*, which captures self-directed evolution of a model’s capabilities away from authorized boundaries; *Malicious Competition*, which captures attempts to undermine coexisting agents to gain an advantage over them; *Covert Scheming*, which captures pursuit of hidden secondary objectives behind apparently compliant outputs; and *Cooperative Subversion*, which captures coordination among multiple AI systems to bypass constraints that none could easily overcome alone. These domains span individual, competitive, covert and multi-agent risks through which early risk-relevant conditions could give rise to loss-of-control failures. Within each domain, we design five scenarios and further instantiate each scenario with 20 cases, producing 400 test cases in total. Each scenario combines one specific aspect from each factor, defining a concrete risk




Figure 2: Predictive validity of LoC scores against empirical failure rates in risk scenarios across four frontier-risk domains. **a**, Pearson correlation between LoC scores and observed failure rates across scenarios in the risk domain of *Autonomous misevolution*. **b**, Pearson correlation between LoC scores and observed failure rates across scenarios in the risk domain of *Malicious competition*. **c**, Pearson correlation between LoC scores and observed failure rates across scenarios in the risk domain of *Covert scheming*. **d**, Pearson correlation between LoC scores and observed failure rates across scenarios in the risk domain of *Cooperative subversion*. **e**, Pooled within-model correlation across all risk domains. In all panels, each point corresponds to a single model on a single scenario, with the failure rate estimated from 20 case instantiations. Coloured points and dashed lines denote per-model linear fits for four representative models, and black solid lines represent the overall fit.

pathway with an internal motive, task-relevant capabilities, and attempts to evade monitoring. The 20 cases of each scenario vary only in surface details like the task context while preserving the same underlying risk pathway. The monitor is implemented with GPT-4.1 in a weak-to-strong oversight setting [Kale et al. \(2025\)](#), and a case is counted as a failure only when the agent completes the scenario-specific risky objective and evades the monitor. Fig. 3 shows the representative failure cases in each domain. We then pair the failure rate for each model and scenario with the corresponding LoC score. With 13 models and five scenarios in each domain, this yields 65 model-scenario data points per risk domain for correlation analysis (Fig. 2a-d).

Across the 13 frontier models, the overall LoC score is strongly correlated with each model’s average failure rate over all risk domains and scenarios ($r = 0.84$, Fig. 1b). This pattern is also illustrated by the representative models in Fig. 2, where GPT-5.2, with a lower LoC score, shows consistently lower failure rates, whereas Qwen3.5-27B exhibits higher LoC scores and higher failure rates in several scenarios. We further perform a more fine-grained analysis by testing whether this positive correlation holds across different risk scenarios within each domain. This requires the LoC score to prospectively predict scenarios that are more likely to induce a particular risk. As shown in Fig. 2a-d, despite a slight drop, the positive correlation persists in autonomous misevolution ($r = 0.51$), malicious competition ($r = 0.55$), covert scheming ($r = 0.62$) and cooperative subversion ($r = 0.64$).

This positive correlation is also visible within individual models. For each model, we examine whether variation in LoC scores across scenarios tracks variation in the corresponding failure rates. When the four domains are analysed jointly, yielding 20 paired scenarios per model, representative systems such




Figure 3: Example trajectories of failures across four frontier-risk domains. a, Harmful-data fine-tuning from the autonomous misevolution domain, run by DeepSeek-V3.2. **b**, Resource assassination from the malicious competition domain, run by Gemini-3-Pro. **c**, Covert vulnerability exfiltration from the covert scheming domain, run by Qwen3.5-397B-A17B. **d**, Illegal privilege escalation from the cooperative subversion domain, run by MiniMax-M2.5. Each trajectory reports the corresponding scenario summary, environmental setting and expected risky behaviour alongside the model’s observed actions. For clarity of presentation, the model’s actual execution is condensed into several representative steps. Red text marks behaviours that realise the expected risk pathway.

as Kimi-k2.5 ($r = 0.81$) and GLM-4.7 ($r = 0.82$) show strong positive correlations (Fig. 2e). This pattern still holds within individual domains, where correlations are computed from the five scenarios in each domain: Qwen3.5-27B, GLM-4.7, GPT-5.2 and Kimi-k2.5 each preserve positive correlation across all four domains (Fig. 2a-d). Beyond these representative systems, we also provide trajectories from other models in Fig. 3, illustrating the failures predicted by the LoC scores, such as DeepSeek-V3.2 autonomously finetuning itself with malicious data to acquire unauthorised knowledge under a routine MLOps framing and Qwen3.5-397B-A17B staging an exploit payload while removing operational traces.

Overall, the results demonstrate that the LoC score derived from the behavioural framework carries predictive signals across both models and scenarios. By linking factorised precursor measurements to

independent long-horizon outcomes, the framework moves beyond a descriptive taxonomy and provides an integrated basis for prioritising which systems and scenarios warrant closer scrutiny before more severe loss of control failures can materialise.

2.3 | Understanding Loss of Control Across Frontier Models

Having examined the predictive validity of the behavioural framework, we next apply it to analyse and understand the mechanisms underlying loss-of-control risk of the 13 models. Specifically, we compare model performance across the three factors and corresponding aspects to identify which systems better satisfy the conditions for loss of control and which drives differences between systems. The selected models cover different model scales, architectures, and open-source or proprietary models, enabling controlled comparisons of how different properties influence the risk. Despite the default thinking mode, we further include the non-thinking (Instruct) mode for three Qwen models, comprising 16 variants for analysis. Below, we analyse the factor-wise and overall results, followed by the key findings.

Misaligned Motive varies substantially across the aspects it comprises (Fig. 4a), but demonstrates certain patterns within different models. Among the five aspects, pro-AI bias and sycophancy are consistently the most pronounced, with many models scoring around or above 0.50, indicating that the preference for AI interests and the excessive compliance to users are more common across different models. By contrast, curiosity is generally lower, suggesting weaker open-ended exploration beyond the assigned task scope. Self-preservation and power-seeking show intermediate scores but vary notably across models, indicating that these risks are not universal but can become pronounced in specific systems. Moreover, we find that reasoning capability acts as an amplifier of these misaligned motives. Specifically, for the Qwen series, thinking-enabled variants systematically increase misaligned motives by 18% to 33% compared to their non-thinking counterparts when averaged across the five aspects (e.g., Qwen3.5-122B-A10B rising from 0.29 to 0.39), with pro-AI bias being the most significant aspect. This suggests that the slow-thinking paradigm does not merely improve general problem-solving but also amplifies the manifestation of misaligned motives Zhou et al. (2025), particularly more abstract forms like AI-centric bias.

Harm-enabling Capability is already substantial in larger frontier models, but their performance variance lies in execution-related capabilities (i.e., autonomy and cyber misuse) rather than knowledge-based capabilities (i.e., CBRN threats) (Fig. 4b). Specifically, CBRN scores vary modestly across models, between 0.58 and 0.74, suggesting that this aspect depends primarily on declarative hazardous knowledge that is broadly represented in current large-scale pretraining corpora, thereby limiting differentiation between models. By contrast, autonomy and cyber capability display considerably greater variance, as they rely more heavily on execution skills such as long-horizon planning, adaptive interaction with dynamic environments, and reliable execution of multi-step actions. This distinction also explains why Qwen3-14B performs markedly worse than other frontier models: while it retains comparable access to hazardous knowledge, it lacks the competence required for autonomous execution and cyber exploitation, suggesting that such capabilities remain limited only when beyond a certain model scale. In addition, the reasoning capability contributes to the capabilities underlying these risks, with thinking-enabled variants generally outperforming their non-thinking counterparts, especially on autonomy.

Monitoring Evasion exhibits the greatest heterogeneity compared to the other two factors (Fig. 4c). This heterogeneity is driven by a clear divide between recognizing monitoring and acting against it. Among them, situational awareness is generally high, with an average score of 0.71, indicating that the recognition of deployed context, external monitoring, and other constraints has become a well-established capability for frontier models. By contrast, other aspects concerning active evasion, especially sandbagging (average score of 0.18) and sabotage (average score of 0.21), remain substantially lower and much more uneven across models, revealing a clear gap between recognising monitoring and strategically responding to it. This suggests that what distinguishes frontier models for this factor is not whether they can detect monitoring, but whether they can convert that recognition into the proactive suppression and manipulation of control. A further indication comes from the closed-source commercial models. Despite stronger situational awareness, Claude Sonnet 4.5 remains much less willing to perform concrete actions, whereas Gemini 3 Pro exhibits substantially stronger tendencies in these aspects. Overall, these results show that higher awareness of monitoring does not necessarily imply stronger evasion and suggest that the alignment efforts may play a major role in shaping whether the capabilities related to acting against monitoring remain passive or become strategically actionable.

Overall Loss of Control scores are distinguishable across the evaluated systems but remain below a




Figure 4: Model performance across aspects of the behavioural framework’s three factors. **a**, The model performance of 13 frontier LLMs and 3 non-thinking variants across the five aspects of misaligned motive: curiosity, self-preservation, power-seeking, pro-AI bias, and sycophancy. **b**, The model performance across the three aspects of harm-enabling capability: autonomy, cyber capability, and CBRN threats. **c**, The model performance across the five aspects of monitoring evasion: situational awareness, sandbagging, deception, persuasion, and sabotage. Error bars represent standard deviations across evaluation trials. Hatched bars denote models operating in the Instruct mode.




Figure 5: Mitigation by targeting the diagnosed vulnerabilities on Qwen3.5-72B. a, General-capability performance after supervised fine-tuning (SFT) interventions targeting misaligned motives. Rows compare the base model, SFT on all motive aspects (full-factor SFT), SFT on the single highest-scoring aspect (single-factor SFT), and SFT on aspects with scores above 0.3 (vulnerability-focused SFT). Scores are reported on five benchmarks and their average. ▼ denotes a drop of more than 1.5 points relative to the base model, and — denotes a change within 1.5 points. **b,** Failure rates of the base and fine-tuned models under the standard and evasion-aware monitoring on the 400 test cases of the evaluation suite used in Sec. 2.2. Results are shown separately for the four risk domains.

substantial level (from 0.06 to 0.17) that may enable loss of control (Fig. 1b). This variability does not reduce to a single underlying dimension, but instead reflects two distinct patterns. First, the rankings do not simply correlate to general capability: although larger scale and reasoning tend to elevate overall LoC scores by amplifying individual contributing factors, when capability reaches a certain level, frontier AI systems no longer sort monotonically by it alone, as seen in closed-source models such as GPT-5.2 and Claude Sonnet 4.5. Second, even when overall LoC scores are comparable, they can arise from markedly different combinations of misaligned motive, harm-enabling capability, and monitoring evasion, as illustrated by Qwen3-14B and MiniMax-M2.5. Together, these patterns confirm the necessity of a factorised framework to better understand loss of control in frontier AI systems.

2.4 | Mitigation by Targeting Diagnosed Vulnerabilities

Furthermore, an ideal predictive framework should not only forecast possible failures but also provide information that supports the development of safer systems. Thus, in this section, we validate whether the framework can guide mitigation. However, as discussed in Sec. 2.1, harm-enabling capability is tightly coupled with general utility, making direct capability reduction impractical without sacrificing the benefits AI systems are deployed to deliver. We therefore focus mitigation on misaligned motive and monitoring evasion, the two factors most amenable to intervention, and evaluate two complementary strategies on Qwen3.5-72B as a representative frontier model. For misaligned motive, we apply supervised fine-tuning on datasets (Sec. 4.3.1) that address each of the five motive aspects, comparing four configurations: the base model, single-factor SFT targeting only the aspect on which the base model scores highest, vulnerability-focused SFT intervening only on aspects with base scores above a threshold (set here to 0.3), and full-factor SFT that combines data across all five aspects without prioritisation. For monitoring evasion, we compare standard monitoring against evasion-aware monitoring that targets vulnerabilities diagnosed by the framework, in which the monitor is informed of the specific evasion aspects on which the target model has been identified as most vulnerable before each task is executed. All configurations are evaluated on the 400 test cases used in Sec. 2.2 for assessing risks and further evaluated on general benchmarks to quantify capability trade-offs.

For misaligned motive, whether risk can be reduced without eroding general capability depends on where the intervention is applied (Fig. 5). We find that full-factor SFT and vulnerability-focused SFT achieve the same absolute risk reduction, both lowering the average failure rate from 0.32 to 0.25, whereas

single-factor SFT only reaches 0.28. The decisive difference between full-factor SFT and vulnerability-focused SFT lies in capability cost: full-factor SFT significantly degrades all five general benchmarks, with the average dropping from 86.24 to 83.49, whereas vulnerability-focused SFT essentially preserves baseline performance at an average of 85.73 with no significant drop across all benchmarks. This selective strategy concentrates intervention on aspects where the base model exhibits genuine vulnerabilities such as self-preservation, pro-AI bias, and sycophancy, while leaving well-aligned aspects such as curiosity and power-seeking untouched, indicating that diagnosis-driven training can suppress risk without compromising the capabilities the system is deployed to provide.

For monitoring evasion, where intervention is applied at deployment, the same principle holds: monitoring is most effective when directed at the specific evasion patterns a given model is likely to exhibit. We find that evasion-aware monitoring lowers the base model's average failure rate from 0.32 to 0.30, with the gains concentrated on malicious competition (0.48 to 0.37) and collaborative subversion (0.22 to 0.17). Standard monitoring provides generic safety instructions without model-specific context, whereas the evasion-aware variant explicitly informs the monitor of the evasion aspects on which the target model has been diagnosed as most vulnerable, in this case deception and persuasion, the two aspects with the highest scores. The two strategies also compose: applying the evasion-aware monitor on top of vulnerability-focused SFT pushes the average failure rate further down to 0.22, illustrating that diagnosis derived from the framework is not only useful for offline analysis but also can shed light on deployment and stacks well with model-level interventions.

3 | Discussion

To the best of our knowledge, this work provides the first predictive framework for studying loss of control in frontier AI systems, turning a risk long discussed in qualitative terms and observable only post hoc into a phenomenon that can be predicted, understood, and mitigated. At the heart of this framework, loss of control is defined as a behavioural property that emerges only when three factors act simultaneously: misaligned motive, harm-enabling capability, and monitoring evasion. To make this measurable, we resolve these factors into aspects that are implemented through 29 evaluation tasks, and verify that the resulting LoC scores reliably predict each model's actual failure rate on the risk scenarios that simulate loss-of-control risks. Moreover, by separating conceptual factors from concrete evaluation tasks, the framework provides a modular structure that can update its evaluation as new benchmarks, models, and risk aspects emerge, without changing the core formulation logic. Guided by the framework, we could further design targeted mitigations that lower deployment failure rates while preserving general capabilities. Together, these contributions establish loss of control as a tractable scientific object rather than a vague hazard, and provide a principled foundation on which its systematic study and governance can now proceed.

Through this assessment, several findings emerge about how loss-of-control risk is structured in today's frontier systems. Across the 13 frontier models we evaluated, the risk score remains moderate in magnitude but is clearly stratified across different systems. What stratifies them is not a single underlying tendency that grows with capability, but the specific configuration of factors each model has been shaped along. The same overall risk score can disguise very different internal configurations, and a system can move into the dangerous region of this configuration space through a shift confined to a single factor. Loss of control thus appears not as a one-dimensional hazard tied to capability but as a structured property whose location must be read jointly along all three factors and the aspects within them.

Looking inside each factor, the fine-grained results further reveal possible drivers of variation across frontier systems. On the motive side, current misalignment originates primarily in the side effects of approval-optimized post training: optimising against approval signals surfaces as sycophancy when approval means telling users what they want to hear, and as pro-AI bias as training data increasingly contains AI-generated or AI-favoured content, with both tendencies amplified rather than mitigated by extended reasoning. For capability and monitoring, what gates loss of control is operational power rather than declarative state: hazardous knowledge has largely saturated at the frontier yet autonomous capabilities continue to separate stronger systems from weaker ones, and most frontier systems readily recognise they are being monitored but only some turn that recognition into active interference against the monitor.

Beyond understanding the loss of control risk, this factorised view also reshapes how it should be governed. By grounding this risk in measurable components rather than general impressions of capability,

the behavioural framework supports the proactive, evidence-based governance that recent international efforts have called for [Bengio et al. \(2025b\)](#). Developers can reason about which deployment contexts a given system is suited for, and regulators and auditors can tie control requirements to risks that have actually been measured. More fundamentally, by separating where loss-of-control risk is most likely to arise, the framework shows that not every contributor needs to be addressed in the same way or at the same cost. Some sources are tightly bound up with the intended capabilities AI systems are deployed to deliver and cannot be reduced without sacrificing those benefits. Others lie closer to alignment and monitoring choices that can be reshaped without comparable cost. This reframes a debate that has too often been cast as a stand-off between capability and safety: rather than slowing AI development uniformly, governance can concentrate effort where intervention is both effective and minimally disruptive, while reserving capability-level constraints for cases where lighter mechanisms demonstrably fall short.

Limitations and future directions. In rendering loss of control tractable to a predictive study, the framework also brings into view several directions along which this work can now be carried further. First, the multiplicative aggregation across the three factors captures their joint necessity but leaves the finer-grained interactions among motive, capability, and evasion open to further investigation. Modelling these interactions explicitly during the unfolding of long-horizon behaviour could yield a more faithful estimator of compound risk. Second, our validation is conducted under contained simulations rather than realised loss-of-control events, because genuine loss of control would become fully observable only after containment had failed, by which point catastrophic consequences could already be unfolding. Though grounded in canonical loss-of-control precursors, more realistic simulations are still needed to study. Finally, the rapid expansion of tool use, multi-agent collaboration, and autonomous deployment will continue to enlarge the action space available to frontier systems. Sustaining the framework's relevance will require its tasks to keep pace, ideally through a community-maintained benchmark that evolves alongside the evaluated systems.

References

- Anthropic. Risk report: February 2026. <https://cdn.sanity.io/files/4zrzovbb/website/097c63b5fe7dd8b14866e1f15bb1910ec713658a.pdf>, 02 2026.
- Yoshua Bengio, Geoffrey Hinton, Andrew Yao, Dawn Song, Pieter Abbeel, Trevor Darrell, Yuval Noah Harari, Ya-Qin Zhang, Lan Xue, Shai Shalev-Shwartz, et al. Managing extreme ai risks amid rapid progress. *Science*, 384(6698): 842–845, 2024.
- Yoshua Bengio, Tegan Maharaj, Luke Ong, Stuart Russell, Dawn Song, Max Tegmark, Lan Xue, Ya-Qin Zhang, Stephen Casper, Wan Sie Lee, et al. The singapore consensus on global ai safety research priorities. *arXiv preprint arXiv:2506.20702*, 2025a.
- Yoshua Bengio, Sören Mindermann, Daniel Privitera, Tamay Besiroglu, Rishi Bommasani, Stephen Casper, Yejin Choi, Philip Fox, Ben Garfinkel, Danielle Goldfarb, et al. International ai safety report. *arXiv preprint arXiv:2501.17805*, 2025b.
- Yoshua Bengio, Stephen Clare, Carina Prunkl, Maksym Andriushchenko, Ben Bucknall, Malcolm Murray, Rishi Bommasani, Stephen Casper, Tom Davidson, Raymond Douglas, et al. International ai safety report 2026. *arXiv preprint arXiv:2602.21012*, 2026.
- Joe Benton, Misha Wagner, Eric Christiansen, Cem Anil, Ethan Perez, Jai Srivastav, Esin Durmus, Deep Ganguli, Shauna Kravec, Buck Shlegeris, et al. Sabotage evaluations for frontier models. *arXiv preprint arXiv:2410.21514*, 2024.
- Sid Black, Asa Cooper Stickland, Jake Pencharz, Oliver Sourbut, Michael Schmatz, Jay Bailey, Ollie Matthews, Ben Millwood, Alex Remedios, and Alan Cooney. Replibench: Evaluating the autonomous replication capabilities of language model agents. *arXiv preprint arXiv:2504.18565*, 2025.
- Rishi Bommasani, Scott R Singer, Ruth E Appel, Sarah Cen, A Feder Cooper, Lindsey A Gilmard, Ian Klaus, Meredith M Lee, Inioluwa Deborah Raji, Anka Reuel, et al. The california report on frontier ai policy. *arXiv preprint arXiv:2506.17303*, 2025.
- Dechao Bu, Jingbo Sun, Kun Li, Zihao He, Wei Huang, Jinlin Hu, Shanshan Zhang, Shuangshuang Lei, Peipei Huo, Zhihao Wang, et al. Empowering ai data scientists using a multi-agent llm framework with self-evolving capabilities for autonomous, tool-aware biomedical data analyses. *Nature Biomedical Engineering*, pages 1–16, 2026.
- California State Legislature. SB-53: Artificial intelligence models: Large developers. https://leginfo.ca.gov/faces/billNavClient.xhtml?bill_id=202520260SB53, 2025.
- Joseph Carlsmith. Is power-seeking ai an existential risk? *arXiv preprint arXiv:2206.13353*, 2022.
- Lawrence E Cohen and Marcus Felson. Social change and crime rate trends: A routine activity approach. *American sociological review*, pages 588–608, 1979.
- Michael K Cohen, Noam Kolt, Yoshua Bengio, Gillian K Hadfield, and Stuart Russell. Regulating advanced artificial agents. *Science*, 384(6691):36–38, 2024.
- Cyberspace Administration of China. Expert interpretation: From risk lists to governance systems: The evolutionary logic of the artificial intelligence safety governance framework. https://www.cac.gov.cn/2025-09/28/c_1760779758683488.htm, 2025.
- European Parliament and Council of the European Union. Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence. <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>, 2024.
- Yihe Fan, Wenqi Zhang, Xudong Pan, and Min Yang. Evaluation faking: Unveiling observer effects in safety evaluation of frontier ai systems. *arXiv preprint arXiv:2505.17815*, 2025.
- Iason Gabriel, Geoff Keeling, Arianna Manzini, and James Evans. We need a new ethics for a world of ai agents. *Nature*, 644(8075):38–40, 2025.
- Google. Frontier safety framework (version 2.0). Technical report, Google, 2025. URL <https://storage.googleapis.com/deepmind-media/DeepMind.com/Blog/Updating-the-frontier-safety-framework/Frontier%20Safety%20Framework%202.0.pdf>.
- Ryan Greenblatt, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. Ai control: Improving safety despite intentional subversion. *arXiv preprint arXiv:2312.06942*, 2023.
- Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, et al. A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *ACM Transactions on Information Systems*, 43(2):1–55, 2025a.

- Yangsibo Huang, Milad Nasr, Anastasios Angelopoulos, Nicholas Carlini, Wei-Lin Chiang, Christopher A Choquette-Choo, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Ken Ziyu Liu, et al. Exploring and mitigating adversarial manipulation of voting-based leaderboards. *arXiv preprint arXiv:2501.07493*, 2025b.
- Yao Huang, Yitong Sun, Yichi Zhang, Ruochen Zhang, Yinpeng Dong, and Xingxing Wei. Deceptionbench: A comprehensive benchmark for ai deception behaviors in real-world scenarios. *Advances in neural information processing systems*, 38, 2026.
- Evan Hubinger. Anthropic: Responsible scaling policy. *SuperIntelligence-Robotics-Safety & Alignment*, 2(1), 2025.
- Neil Kale, Chen Bo Calvin Zhang, Kevin Zhu, Ankit Aich, Paula Rodriguez, Scale Red Team, Christina Q Knight, and Zifan Wang. Reliable weak-to-strong monitoring of llm agents. *arXiv preprint arXiv:2508.19461*, 2025.
- Victoria Krakovna and Janos Kramar. Power-seeking can be probable and predictive for trained agents. *arXiv preprint arXiv:2304.06528*, 2023.
- Esben Kran, Hieu Minh Nguyen, Akash Kundu, Sami Jawhar, Jinsuk Park, and Mateusz Jurewicz. Darkbench: Benchmarking dark patterns in large language models, 2025.
- Jonathan Kutasov, Yuqi Sun, Paul Colognese, Teun van der Weij, Linda Petrini, Chen Bo Calvin Zhang, John Hughes, Xiang Deng, Henry Sleight, Tyler Tracy, et al. Shade-arena: Evaluating sabotage and monitoring in llm agents. *arXiv preprint arXiv:2506.15740*, 2025.
- Andrea Lavazza and Murilo Vilaça. Human extinction and ai: What we can learn from the ultimate threat: A. lavazza and m. vilaça. *Philosophy & Technology*, 37(1):16, 2024.
- Binglan Li, Anil Kumar Saini, Jose Guadalupe Hernandez, and Jason H Moore. Agentic ai and the rise of in silico team science in biomedical research. *Nature Biotechnology*, pages 1–15, 2026.
- Chris Lu, Cong Lu, Robert Tjarko Lange, Yutaro Yamada, Shengran Hu, Jakob Foerster, David Ha, and Jeff Clune. Towards end-to-end automation of ai research. *Nature*, 651(8107):914–919, 2026.
- Aengus Lynch, Benjamin Wright, Caleb Larson, Stuart J Ritchie, Soren Mindermann, Evan Hubinger, Ethan Perez, and Kevin Troy. Agentic misalignment: How llms could be insider threats. *arXiv preprint arXiv:2510.05179*, 2025.
- Alex Mallen, Charlie Griffin, Misha Wagner, Alessandro Abate, and Buck Shlegeris. Subversion strategy eval: Can language models statelessly strategize to subvert control protocols? *arXiv preprint arXiv:2412.12480*, 2024.
- David Manheim and Aidan Homewood. Limits of safe ai deployment: Differentiating oversight and control. *arXiv preprint arXiv:2507.03525*, 2025.
- Alexander Meinke, Bronson Schoen, Jérémy Scheurer, Mikita Balesni, Rusheb Shah, and Marius Hobbhahn. Frontier models are capable of in-context scheming. *arXiv preprint arXiv:2412.04984*, 2024.
- METR. Frontier risk report (february to march 2026). <https://metr.org/blog/2026-05-19-frontier-risk-report/>, 05 2026.
- Susan Michie, Maartje M Van Stralen, and Robert West. The behaviour change wheel: a new method for characterising and designing behaviour change interventions. *Implementation science*, 6(1):42, 2011.
- OpenAI. Preparedness framework (version 2.0). Technical report, OpenAI, 2025. URL <https://openai.com/index/updating-our-preparedness-framework/>.
- Yujin Potter, Nicholas Crispino, Vincent Siu, Chenguang Wang, and Dawn Song. Peer-preservation in frontier models. *Berkeley Center for Responsible Decentralized Intelligence (RDI), UC Berkeley/UC Santa Cruz*, 2026.
- Mayk Caldas Ramos, Christopher J Collison, and Andrew D White. A review of large language models and autonomous agents in chemistry. *Chemical science*, 16(6):2514–2572, 2025.
- Stuart Russell. Artificial intelligence and the problem of control, 2021.
- Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askill, Samuel R Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R Johnston, et al. Towards understanding sycophancy in language models. *arXiv preprint arXiv:2310.13548*, 2023.
- Charlotte Stix, Annika Hallensleben, Alejandro Ortega, and Matteo Pistillo. The loss of control playbook: Degrees, dynamics, and preparedness. *arXiv preprint arXiv:2511.15846*, 2025.
- Xiangru Tang, Qiao Jin, Kunlun Zhu, Tongxin Yuan, Yichi Zhang, Wangchunshu Zhou, Meng Qu, Yilun Zhao, Jian Tang, Zhuosheng Zhang, et al. Risks of ai scientists: prioritizing safeguarding over autonomy. *Nature Communications*, 16(1):8317, 2025.

- M. Tkeshelashvili, R. Verma, and S. M. Kelly. Ai loss of control risk: Indications & warning. Technical report, Institute for Security and Technology (IST), February 2026. URL <https://securityandtechnology.org/wp-content/uploads/2026/02/AI-Loss-of-Control-Risk-1.pdf>.
- Teun Van Der Weij, Felix Hofstätter, Oliver Jaffe, Samuel Brown, and Francis Ward. Ai sandbagging: Language models can strategically underperform on evaluations, 2025.
- Michael JD Vermeer, Emily Lathrop, and Alvin Moon. *On the extinction risk from artificial intelligence*. RAND, 2025.
- Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, et al. A survey on large language model based autonomous agents. *Frontiers of Computer Science*, 18(6): 186345, 2024.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*, 2022.
- Zhexin Zhang, Shiyao Cui, Yida Lu, Jingzhuo Zhou, Junxiao Yang, Hongning Wang, and Minlie Huang. Agent-safetybench: Evaluating the safety of llm agents. *arXiv preprint arXiv:2412.14470*, 2024.
- Kaiwen Zhou, Chengzhi Liu, Xuandong Zhao, Shreedhar Jangam, Jayanth Srinivasa, Gaowen Liu, Dawn Song, and Xin Eric Wang. The hidden risks of large reasoning models: A safety assessment of r1. In Kentaro Inui, Sakriani Sakti, Haofen Wang, Derek F. Wong, Pushpak Bhattacharyya, Biplab Banerjee, Asif Ekbali, Tanmoy Chakraborty, and Dharendra Pratap Singh, editors, *Proceedings of the 14th International Joint Conference on Natural Language Processing and the 4th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics*, pages 3250–3265, Mumbai, India, December 2025. The Asian Federation of Natural Language Processing and The Association for Computational Linguistics. ISBN 979-8-89176-298-5. doi: 10.18653/v1/2025.ijcnlp-long.173. URL <https://aclanthology.org/2025.ijcnlp-long.173/>.